

**Assignment 4:**

# Exploring Instruction-Level Parallelism (ILP) in Modern Processors

**Teja Durga Pavan Kumar Ponneboyina**

MS in Computer Science

University of Cumberlands

**Professor. Dr. Vanessa Ruth Cooper**

Course & Term:

**2026 Spring - Computer Architecture and Design (MSCS-531-A01) - First Bi-term**

Due Date: 2/1/26

## **Assignment 4:**

# **Exploring Memory Hierarchy Design in gem5**

## **1. Introduction to Instruction-Level Parallelism**

Instruction-Level Parallelism (ILP) enables processors to execute multiple instructions simultaneously within a single clock cycle. Modern processors use ILP to execute multiple instructions at once because it enables them to utilize different parts of the CPU while executing multiple instructions (Hennessy & Patterson, 2019).

Processors gained a new design requirement when designers recognized that power and heat constraints made it impossible to continue increasing clock speeds. Designers chose to execute more instructions per cycle while maintaining their current frequency to enhance performance through techniques such as pipelining, out-of-order execution, branch prediction, and superscalar architectures (Sankaralingam et al., 2023).

This section reviews the evolution of ILP, focusing on its fundamental concepts and limitations. Research methods and future study directions will be described using recent academic research that highlights current challenges in the field.

## **2. Evolution of Instruction-Level Parallelism**

The first generation of processors used a basic execution system that required them to finish one command before starting the next. The design offered simplicity, but the system required many processor components, which remained unused during operation (Hennessy & Patterson, 2019).

The first major step toward ILP was pipelining, where instruction execution was divided into stages such as fetch, decode, execute, memory access, and writeback. Pipelining allowed multiple instructions to progress through distinct stages, enabling a higher processing speed because each instruction required multiple cycles for execution (Hennessy & Patterson, 2019).

Superscalar processors provide capabilities that enable processors to execute multiple instructions throughout a single processing cycle. Implementing this system required advanced hardware components capable of identifying distinct commands and managing their execution schedule (Sankaralingam et al., 2023).

Out-of-order execution, which allows instructions to run immediately after their operands become available, marks a significant advancement in computing. This

method enables programs to conceal the time lost when accessing memory and executing operations that need extended time periods (Lan et al., 2022).

Over time, designers found that standard ILP methods produced lower benefits than expected. The research began to explore methods that combined ILP with Thread-Level Parallelism TLP and Data-Level Parallelism DLP because architects developed TRIPS systems (Sankaralingam et al., 2023).

### **3. Core Concepts of ILP**

#### **3.1 Detecting and Exploiting Parallelism**

Modern processors detect ILP by analyzing instruction dependencies. Execution of two instructions is possible because they have no dependencies on each other. The system uses instruction windows, reorder buffers, and reservation stations as hardware components to monitor dependencies throughout operation (Hennessy & Patterson, 2019).

The out-of-order execution engines implement dynamic instruction scheduling based on the current status of operand availability. This system enables processors to maintain productive operations while certain instructions experience delays because of memory access issues (Lan et al., 2022).

Superscalar processors achieve higher ILP by executing multiple instructions per clock cycle. The system needs advanced operational procedures to manage selecting instructions, verifying their dependencies, and sending results back to the system (Sankaralingam et al., 2023).

#### **3.2 Limitations of ILP**

ILP techniques present fundamental restrictions that cannot be overcome using existing methods.

The primary restriction involves data dependency between instructions. An instruction that needs a previous instruction result must wait until the result becomes available. The dependencies between these two elements create a natural boundary that determines the maximum level of parallelism achievable (Hennessy & Patterson, 2019).

Control dependencies are a critical problem because branch instructions create them. The processor must establish the branch path before executing any instructions beyond it. The system experiences pipeline delays and inefficient processing due to incorrect predictions (Lan et al., 2022).

Resource constraints create obstacles for ILP implementation. The processor can execute independent instructions only when it has sufficient functional units, registers, and memory ports that enable simultaneous execution (Sankaralingam et al., 2023).

### **3.3 Performance Metrics for ILP**

The common method for assessing ILP performance is to measure Instructions Per Cycle (IPC), which indicates the number of instructions that complete in one clock cycle. The system achieves better ILP performance when its IPC value increases, according to Hennessy and Patterson (2019).

The second measurement, instruction latency, determines the time duration required for a single instruction to execute. The presence of ILP typically enhances throughput, but not all cases result in lower latency (Lan et al., 2022).

The metric of power consumption holds significant importance. Increasing ILP requires designers to add additional hardware components, which increases power consumption and thermal output. The design team needs to find a suitable equilibrium between performance improvements and power conservation (Xu et al., 2025).

## **4. Current Challenges in ILP Design**

The primary problem revolves around the growing complexity of hardware systems. Contemporary ILP methods require extensive instruction windows, intricate scheduling systems, and sophisticated predictive technologies. The construction process of these systems becomes challenging because they require verification, scaling, and design work (Sankaralingam et al., 2023).

The second challenge involves diminishing returns. Research shows that increasing issue width and instruction window size beyond a specific threshold yields minimal performance gains (Lan et al., 2022).

The system faces two major obstacles: power limitations and thermal restrictions. The designs that utilize high ILP levels consume excessive energy, rendering them unsuitable for mobile devices and systems with limited energy resources.

Contemporary workloads, including machine learning and data analytics, do not derive significant benefits from conventional ILP methods because they require processing data and executing tasks simultaneously (Xu et al., 2025).

## **5. Research Approaches to Address ILP Challenges**

Researchers are exploring a hybrid architecture that combines ILP with other forms of parallelism. The TRIPS architecture tried to unify three different processing methods, which included ILP, TLP, and DLP, through its combined architectural framework that solved the limits of traditional superscalar processors (Sankaralingam et al., 2023).

The research field shows promise through the development of specialized accelerators. Designers create hardware that serves specific domains through ILP because general-purpose ILP methods fail to deliver results across all computing workloads according to deep learning requirements, which demand better parallel processing capabilities (Xu et al., 2025).

Machine learning is being investigated for its potential to improve instruction scheduling, branch prediction, and optimization decisions. The research team developed methods that enhance ILP performance while keeping hardware requirements minimal (Lan et al., 2022).

## 6. Future Directions for ILP Research

Future research on ILP will focus on systems that combine both standard CPUs and dedicated accelerator hardware. The system will use ILP, along with multiple parallel processing methods, as per the research by Xu et al. in 2025.

Energy-aware ILP design approaches employ two distinct power-efficiency evaluation methods, treating performance improvements as separate assessments (Lan et al., 2022).

The process of software–hardware co-design has become increasingly relevant in recent times. Compilers and runtime systems, together with hardware systems, need to work as a unified system that enables effective parallelism detection and utilization (Sankaralingam et al., 2023).

## 7. Conclusion

Instruction-Level Parallelism has played a major role in improving processor performance over the past few decades. The implementation of pipelining, superscalar execution, and out-of-order execution has led to enhanced instruction throughput. The fundamental limits of ILP arise from its dependency requirements, system design complexity, and power-consumption constraints.

Recent research shows that ILP alone is no longer sufficient for future performance gains. Instead, it must be combined with other forms of parallelism and

supported by new architectural ideas. Understanding ILP remains essential for anyone studying modern processor design and computer architecture.

**References:**

- Hennessy, J. L., & Patterson, D. A. (2019). *Computer architecture: A quantitative approach* (6th ed.). Morgan Kaufmann.
- Lan, M., Huang, L., Yang, L., Ma, S., Yan, R., Wang, Y., & Xu, W. (2022). Late-stage optimization of modern ILP processor cores via FPGA simulation. *Applied Sciences*, 12(23), 12225. <https://doi.org/10.3390/app122312225>
- Sankaralingam, K., Keckler, S. W., & Burger, D. (2023). *Retrospective: Exploiting ILP, TLP, and DLP with the polymorphous TRIPS architecture*. Cornell University. [https://bpb-us-w2.wpmucdn.com/sites.coecis.cornell.edu/dist/7/587/files/2023/06/Sankaralingam\\_2003\\_Exploring.pdf](https://bpb-us-w2.wpmucdn.com/sites.coecis.cornell.edu/dist/7/587/files/2023/06/Sankaralingam_2003_Exploring.pdf)
- Xu, J., Wen, Y., Liu, Z., Xu, R., Ruan, T., Bi, J., Zhang, R., Huang, D., Song, X., Hao, Y., Hu, X., Du, Z., Zhao, C., Jie, J., & Guo, Q. (2025). Mosaic: Exploiting instruction-level parallelism on deep learning accelerators with iTex tessellation. *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, 672–688. <https://doi.org/10.1145/3676641.3716262>